

Condition–action rules in controlling complex systems

Sotiris Moschoyiannis

joint work with

Matthew R. Karlsen and Vlad Georgiev

University of Surrey, England

RuleML Webinar @ Skype

29th March 2019

Outline

Introduction

Condition-action rules for learning individual passengers' preferences on transport networks

- XCSI and Journey Recommendations

- Experiments and Simulations

- Results and Discussion

Condition-action rules in controlling RBNs

- Random Boolean Networks (RBNs)

- XCS

Condition-action rules for controlling dynamical systems

- Systems dynamics models

- XCSR

Concluding Remarks

Introduction

OJPA Project

Onward Journey Planning Assistant

Personalised recommendations for:

- ▶ customers planning a journey
- ▶ customers experiencing disruption

Overall challenge

- ▶ Input:

- ▶ environment factors:

- ▶ train, taxi, tube, boat and bus [0 or 5]; weather [0 to 5]

- ▶ journey-specific factors:

- ▶ current delay, delay on current mode(s), onward delay [0 or 5]

- ▶ passenger preferences:

- ▶ value, speed, comfort, shelter [0 to 5]

- ▶ Output:

- ▶ “Correct” single integer recommendation:

- ▶ no change (0), single taxi (1), shared taxi (2), bus (3), boat (4), tube (5), train (6)

Population of Rules (the Knowledge)

Condition	: Action
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [#,#] [2,5] [#,#] [#,#]	: 1
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [2,5] [#,#] [#,#] [#,#]	: 1
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [#,#] [4,5] [#,#] [#,#]	: 1
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [4,5] [#,#] [#,#] [#,#]	: 1
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [0,1] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [2,3] [0,1] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [2,3] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [0,1] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [2,3] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [0,1] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [0,1] [4,5] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [2,3] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [2,3] [4,5] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [4,5] [0,1] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [4,5] [2,3] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [4,5] [4,5] [#,#] [#,#]	: 2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [0,1] [0,1] [#,#] [#,#]	: 3
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [0,1] [0,1] [#,#] [#,#]	: 3
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [2,3] [0,1] [#,#] [#,#]	: 3
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [0,1] [0,1] [#,#] [#,#]	: 3
...	: ...

Rule Matching

Figure 1: Simple rule matching example

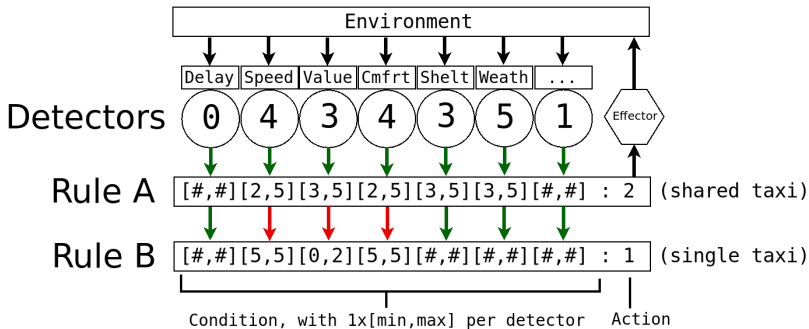
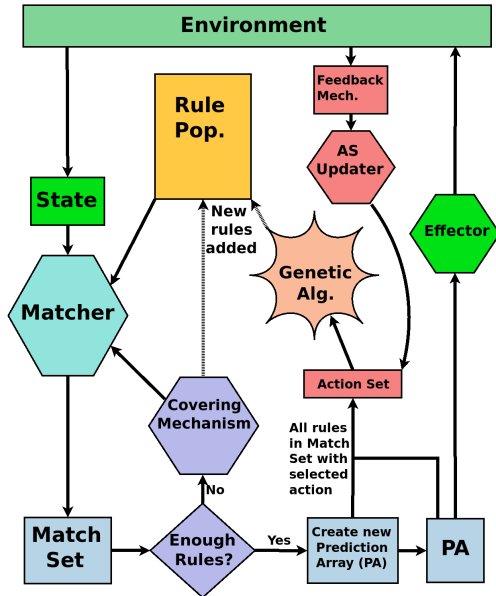


Figure 2: XCSI overview



Effector and Feedback Mechanism

- ▶ Effector outputs:
 - ▶ no change (0), single taxi (1), shared taxi (2), bus (3), boat (4), tube (5), train (6)
- ▶ Feedback Mechanism, receives:
 - ▶ 1000 for a correct suggestion
 - ▶ 0 for an incorrect suggestion
- ▶ This feedback is used to update the action set rules

XCS details in: Butz, Martin V., and Stewart W. Wilson. "An algorithmic description of XCS." International Workshop on Learning Classifier Systems. Springer, Berlin, Heidelberg, 2000.

Experiments

- ▶ Simulation based on London tube network
- ▶ 300 artificial 'passengers' with randomised:
 - ▶ preferences [0 to 5]
 - ▶ origin location / station
 - ▶ destination location / station
- ▶ Each passenger takes multiple journeys

Simulation details (1)

- ▶ Random starting time step (0 to 99) for each passenger
- ▶ Each time step has weather [0 to 5; random]
- ▶ Train, boat, bus and taxi [0 or 5; random]
- ▶ Passenger is at one node each time step
- ▶ In each time step 5% of links are out-of-action
- ▶ Interleaved... (see next 3 slides)

Simulation details (2)

Figure 3: Input list production, step 1 – order

Time Step 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 ...

Journey 1	S	S	S	S	S	S										
Journey 2		S	S	S	S	S	S	S	S							
Journey 3					S	S	S	S	S	S	S	S	S	S	S	S
Journey 4	S	S	S	S	S	S	S	S								
Journey 5			S	S	S	S	S									
Journey 6							S	S	S	S	S	S	S	S	S	
Journey 7	S	S	S	S	S											

S = 'journey state'

Simulation details (3)

Figure 4: Input list production, step 2 – shuffle

Time Step 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 ...

Journey 3					S	S	S	S	S	S	S	S	S	S	S	S
Journey 1	S	S	S	S	S	S										
Journey 7	S	S	S	S	S											
Journey 6							S	S	S	S	S	S	S	S	S	
Journey 4	S	S	S	S	S	S	S	S								
Journey 2		S	S	S	S	S	S	S	S							
Journey 5			S	S	S	S	S									

S = 'journey state'

Simulation details (4)

Figure 5: Input list production, step 3 – obtain input vectors

Time Step 0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 ...

Journey 3					S	S	S	S	S	S	S	S	S	S	S	S
Journey 1	I1	I4	I8	S	S	S										
Journey 7	I2	I5	I9	S	S											
Journey 6							S	S	S	S	S	S	S	S	S	
Journey 4	I3	I6	...	S	S	S	S	S								
Journey 2		I7	S	S	S	S	S	S	S							
Journey 5			S	S	S	S	S									

S = 'journey state'

IX = 'Input X'

Simulation details (5)

Input is checked against the 'real world' preferences of the simulated customers to get a correct input output pair...


Condition	:	Action
(0)(0)(0)(5)(5)(3)(2)(5)(1)(0)(3)(4)(4)	:	?
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [2,3] [4,5] [#,#] [#,#] [#,#]	:	1
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [0,1] [0,1] [2,3] [#,#] [#,#]	:	2
[0,0] [0,0] [0,0] [5,5] [5,5] [#,#] [#,#] [5,5] [4,5] [0,1] [0,1] [#,#] [#,#]	:	3
...	:	...

Correct input–output pair in this example: 0005532510344:2

We therefore assemble a list of inputs and answers (> 51,000) for training and testing XCSI.


Results (1)

Figure 6: Error % for different parameter settings (1 of 2)



Results (2)

Figure 7: Error % for different parameter settings (2 of 2)



Final Parameters (Adjusted Only)

Parameter	Value	Brief Description
N	11700	Rule population size
$P_{\#}$	0.35	Probability of hash
ϵ_0	5	Error threshold
θ_{ga}	25	Genetic algorithm frequency
θ_{del}	10	Deletion threshold
β	0.1	Affects update of p, ϵ , and action set size for classifiers
α	0.08	Affects fitness updates
ν	6	Affects fitness updates
χ	1	Likelihood of GA crossover operation
μ	0.08	Likelihood of GA mutation operation
δ	0.05	Modifies the effect of fitness on classifier 'deletion vote'
θ_{sub}	60	Subsumption threshold
AS subsumpt.	false	Perform subsumption in the action set?

Final Performance

- ▶ Minimum trial error: 0.1%
- ▶ Average error: 0.734%
- ▶ Maximum trial error: 2%
- ▶ In concrete terms, over 99% of passengers would get the correct suggestion.

Outline

Introduction

Condition-action rules for learning individual passengers' preferences on transport networks

- XCSI and Journey Recommendations

- Experiments and Simulations

- Results and Discussion

Condition-action rules in controlling RBNs

- Random Boolean Networks (RBNs)

- XCS

Condition-action rules for controlling dynamical systems


- Systems dynamics models

- XCSR

Concluding Remarks


Random Boolean Networks (RBNs)

Figure 8: A Random Boolean Network (RBN) with $N=3$, $K=2$




Controllability in RBNs (1/3)

Figure 9: State space of RBN of Fig. 8 (all AND); two attractors



Controllability in RBNs (2/3)

Figure 10: State space of RBN of Fig. 8 (all XOR); two attractors



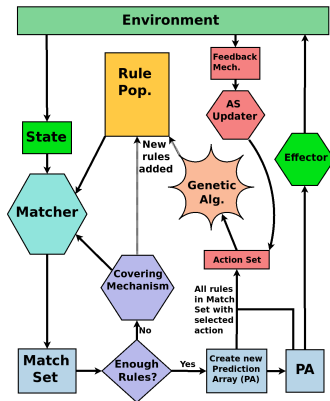
Controllability in RBNs (3/3)

- ▶ The general notion of *control* in complex networks:
 - ▶ The ability to direct the network from any state to a target state.
- ▶ can be refined to:
 - ▶ the ability to direct an RBN from any state to (one of) its attractors

The objective is to evolve a rule set that directs an RBN from any state to (one of) its attractors.

XCS – overview

Figure 11: XCS - **condition** is a bit string





XCS details in: Martin V. Butz and Stewart W. Wilson. "An algorithmic description of XCS." Int'l Workshop on *Learning Classifier Systems*. Springer, Berlin, 2000.

Applying XCS to control RBNs

- ▶ Each rule represents a *condition : action* expression that links specific states of the RBN (conditions) to bit flips (actions)
- ▶ To shift from single state to the state cycle attractor, apply one of ###:1; ###:2; ###:3
- ▶ To shift from the state cycle to the single state attractor, apply one of 110:3; 011:1; 101:2; 001:3; 010:2; 100:1
 - ▶ where # denotes "don't care" and the action represents the index of the bit to flip

XCS, not XCSI

Figure 12: Rules in XCS



Controlling RBNs using XCS

Figure 13: Control graph for a $N=5, K=2$ RBN using XCS

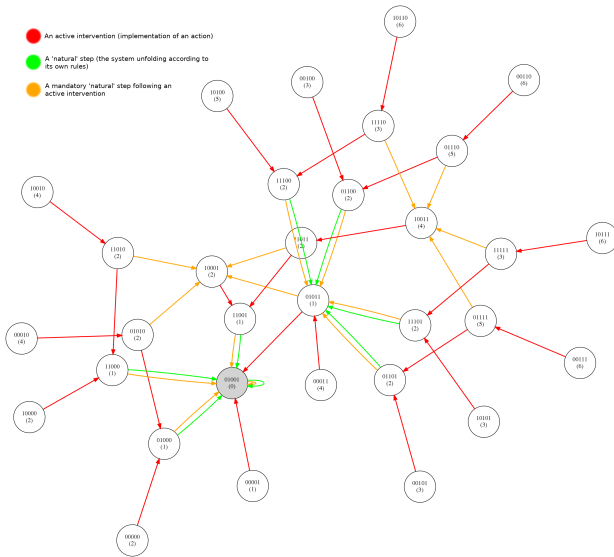


Figure for OCR vs XCS

*Wilson, Stewart W. "Mining oblique data with XCS." International Workshop on Learning Classifier Systems. Springer, Berlin, Heidelberg, 2000.

Outline

Introduction

Condition-action rules for learning individual passengers' preferences on transport networks

- XCSI and Journey Recommendations

- Experiments and Simulations

- Results and Discussion

Condition-action rules in controlling RBNs

- Random Boolean Networks (RBNs)

- XCS

Condition-action rules for controlling dynamical systems

- Systems dynamics models

- XCSR

Concluding Remarks

System Dynamics Models

- ▶ System dynamics models¹ are often understood as complex networks
- ▶ *Control* involves interventions on the values of state variables (*actions*) to steer the network to a desired state
- ▶ XCSR (*XCS Real*): real valued extension of XCS

Question: Can we use XCSR to identify correct actions on *control* nodes² for steering a dynamical system to a desired state ?

¹J.D. Sterman "Business Dynamics: Systems Thinking and Modelling for a Complex World", McGraw-Hill, New York, 2000

²S. Moschoyiannis, N. Elia, A. Penn, *et al* "A Web-based Tool for Identifying Strategic Intervention Points in Complex Systems ", ETAPS 2016

Population growth model (1/2)

The S-shaped population growth model³:

$$\frac{d\text{Population}}{dt} = \text{Birth Rate} - \text{Death Rate} \quad (1)$$

$$\text{Birth Rate} = \text{FractionalBR} * \text{Population} \quad (2)$$

$$\text{Death Rate} = \text{FractionalDR} * \text{Population} \quad (3)$$

$$\text{FractionalBR} = 1 - \frac{1}{1 + e^{-7 * (\text{PCC} - 1)}} \quad (4)$$

$$\text{PCC} = \frac{\text{Population}}{\text{Carrying Capacity}} \quad (5)$$


and the value of Carrying Capacity is a constant.

³J.D. Sterman "Business Dynamics: Systems Thinking and Modelling for a Complex World", McGraw-Hill, New York, 2000

Population growth model (2/2)

- System has only one state variable, thus one attractor where the rates converge

Figure 14: Population (blue) converges at the value of Carrying Capacity (grey); BirthRate(red) and DeathRate(yellow) converge at half that value



Controlling this model (1/2)

- ▶ Establish *continuous* control
 - ▶ *Step* defined as the change of value in state variable after one unit of time
- ▶ XCSR selects an action before each step
 - ▶ Action set: *increase/decrease* Birth- and Death-Rate by a constant or percentage, or no action for a given step
- ▶ Set *explore / exploit* to allow for 50% chance of selecting random action

Controlling this model (2/2)

XCSR can bring the system to a desired stable state (attractor)

- ▶ **But** the evolved rule set contains *overgeneralised* rules (too many #'s)
- ▶ **Revisited** reward mechanism – generalisation from mutation only
- ▶ **But** when current state is "too far" from desired state => *action chain learning problem*
- ▶ **Revisited** notion of *step*: action selection followed by full simulation until convergence

Lotka-Volterra predator-prey model (1/2)

This model⁴ is a well known non-linear first order differential equations couple:

$$\frac{d\text{Prey}}{dt} = \text{Prey} * (\text{PreyBirthRate} - \text{DeathRateperPredator} * \text{Predator}) \quad (6)$$

$$\frac{d\text{Predator}}{dt} = -\text{Predator}(\text{PredatorDeathRate} - \text{BirthRateperPrey} * \text{Prey}) \quad (7)$$


where PreyBirthRate , $\text{DeathRatePerPredator}$, PredatorDeathRate and BirthRatePerPrey are constants.

⁴A.J. Lotka "Elements of Mathematical Biology" Dover Publications, New York, 1956

Lotka-Volterra predator-prey model (2/2)

- System has two state variables, thus more than one attractors

Figure 15: Two state variables imply multiple attractors; Predator(red) and Prey(blue) populations have different attractors



Controlling the L-V model (1/2)

XCSR generates logically correct classifiers, and can bring the system to a desired stable state (attractor)

- ▶ **But** when more than 10 consecutive good actions are needed
 - ▶ XCSR not able to reach a reward in the early runs
 - ▶ overgeneralised rules with prediction of 0 reward
 - ▶ hence, these rules are highly accurate, hence their fitness inflated
 - ▶ leads to *action chain learning problem*
- ▶ **Revisited** notion of action – using adaptive step sizes seems promising; *adaptive action mapping (XCSAM)*

Conclusions

- ▶ XCSR can control system dynamics models, under certain conditions
 - ▶ challenge: **multi-step** problems, **continuous** control
- ▶ To address the *action chain learning problem*, convert the multi-step problem to multiple single step problems
- ▶ **However**, not feasible for higher order systems
 - ▶ Try adaptive step sizes \Rightarrow XCSAM
 - ▶ Try discretising the system state
- ▶ **Rule-based machine learning** (RBML) is human readable
 - ▶ why XCS* suggested a given action
 - ▶ how XCS* arrived at suggesting that action

Thanks/Acknowledgements

- ▶ Thank you for your attention
- ▶ Thanks to Alastair Finlinson, Sophia Manalo, George Papagiannis
- ▶ This research was partly funded by
 - ▶ the Department for Transport, via Innovate UK and the *Accelerating Innovation in Rail* (AIR) Round 4 programme
 - ▶ the UKRI EPSRC council, and the AGELink project
 - ▶ the NCSC, and two research summer internships
 - ▶ the EIT Digital, and the Real-Time Flow project